

CS730 Assignment 1: PCI Device Driver

180301 - hrugved wath

March 1, 2022

1 Design

User programs can access the services of cryptocard device using the `cryter.c` library which interacts with the device driver `cryptocard_mod.c` using a combination of sysfs variables and character device.

- Driver - `cryptocard_mod.c`
 - Configuration can be set using following sysfs variables - INTERRUPT(0/1), DMA(0/1), KEY_A(0-255), KEY_B(0-255), DECRYPT(0/1), IS_MAPPED(0/1).
 - DMA=1 and IS_MAPPED=1 is not supported.
 - encryption and decryption requires writing input to character device and reading from the same for getting result back.
 - Limited concurrency support with effectively exposing a lock using a sysfs variable TID. Each task has to lock the driver, set desired configuration, do the processing and unlock the driver.
- Library - `cryter.c`
 - provides a minimal interface for user programs
 - Support encryption/decryption of arbitrary sizes by operating on data in chunks.
 - Provides a multi-device handle support. The device configuration is per-handle, i.e changing any configuration including keys of one handle will not affect other device handles.

2 Implementation

- Driver - `cryptocard_mod.c`
 - All sysfs variables except TID, are in an attribute group controlled by `sysfs_store()` and `sysfs_show()`.
 - `sysfs_store_tid()` and `sysfs_show_tid()` controls sysfs variable TID. TID at any time is either the tid of the task currently operating the device or -1 if free. `sysfs_store()` is protected by a mutex `dev_lock`. Each task tries to acquire `dev_lock`, then writes its tid to TID. Now the task is free to operate on the device. Once done, this task has to store -1 to TID which is accompanied by releasing `dev_lock`.
 - `cdev_write()` is character device write function. If memory is not mapped, it first copies data to `mmio_buf/dma_buf` and then calls `mmio()/dma()`. This function returns only when result is available.
 - `cdev_read()` is called for reading result. The result is guaranteed to be available when called.
 - `cdev_mmap()` is called for mmap the device memory.
 - `write_to_device()` and `read_to_device()` are used for copying kernel buffer to and from device memory respectively.

- `mmio()` and `dma()` either polls the status register in case of non-interrupt mode, else in interrupt mode, it waits on a wait queue (`wq_is_data_ready`) in intereruptable mode.
- `irq_handler()` is called in interrupt mode. It signals the `wq_is_data_ready` wait queue.
- Library - `cryter.c`
 - It maintains a linked list of opened cdevs in head of struct `dev_handle`. Calling `create_handle()` adds a new entry to this list. `dev_handle` apart from cdev, contains device configuration.
 - The linked list has utility functions for adding, deleting and most importantly fetching `dev_handle` of a cdev. All operations are protected by mutex `mutex_list` for multi-threaded support.
 - The `set_*` functions is used by user to set config related to cdev, i.e updating the corresponding `dev_handle`. The `_set_device_config()` is called for actually setting the device configuration. This functions internally uses `_set_device_*` functions to set the individual sysfs variables.
 - `_lock_device()` tries to get hold of driver by writing its tid to sysfs `TID`. This is blocking call as it requires locking the mutex inside driver. `_unlock_device()` writes -1 to `TID`.
 - encrpyt first tries to get hold of driver using `_lock_device()`, it then setup the device config of provided cdev using a `_set_device_config()`. Calls `_device_operate()` repeadeltly with chunks of data. Finally calls `_unlock_device()` to release the device. decrpyt operates in similar fashion.
 - `MMIO_CHUNK_SIZE` and `DMA_CHUNK_SIZE` determines the chunk size for mmio and dma operation respectively. It is the max size of data device can operate on in given mode.
 - `map_card()` and `unmap_card()` shifts the pointer with `DEVICE_MEMORY_OFFSET` which basically points to start of unused memory region where device expects the data.

3 Testing

- Passed all the tests given in eval-test suite.
- systematic testing targeted for checking a particular feature. Like creating multiple device handles in multi-process and multi-threaded environments to check for correct list implementation. Using data of various sizes for both checking both operating in chunk feature and mmap calls.
- Multi-threaded testing with around 10 threads using the device concurrently.
- Ad-hoc testing, trying to find edge cases.
- perturbation to find possible scenarios of deadlocking.

4 Benchmarks

`pidstat 1 -e ./program` : collect stats every 1 second and averages them.

program	%usr	%system	%guest	%wait	%CPU
mmio	0.87	76.80	0.00	22.36	77.67
mmio_interrupt	0.25	7.57	0.00	14.85	7.82
dma	0.02	98.47	0.00	1.50	98.49
dma_interrupt	0.01	0.34	0.00	0.08	0.35
mmap	1.46	78.87	0.00	19.56	80.33
mmap_interrupt	0.05	1.12	0.00	2.09	1.18